

# VQA

Mar 18, 2021 to Mar 24, 2021



# Status Update

---

- We completed case 3 - trial 2, case 2 trial 2.
- We partially completed case 4 trial 2 from experiment indexes 4300-4417.
- Resource check was performed locally for each of the case before execution on the AWS instances.

# Status Update

---

- We stopped the AWS yesterday late evening.
- Training completed for **1078** experiments on AWS alone!
- 304 trainings completed over combinations of MIT DGX1, FSS DGX1 and OpenMind.
- FSS DGX1 will still be available till Mar 26, 2021.

# AWS Update



Over 16,500 hours of computation utilized.

EC2	\$12,439.69
DataTransfer	\$35.21
kms	\$0.00
CloudTrail	\$0.00
Other Services	\$0.00
Tax	\$1,247.48
<b>Total</b>	<b>\$13,722.38</b>

This amount is only for March 2021.  
Feb 2021 cost is \$5.

→ We are currently running 'sep res' in multitask experiments in case 1 trial 1.

Progress Table

Tag	Status	Progress	BestValAcc	Elapsed Time
train_1404	In Progress...	50.8%	0.3505859375	2 days, 15:56:11.408568
train_1405	Finished	100%	0.2724609375	2 days, 9:58:05.993283
train_1406	In Progress...	52.3%	0.57421875	2 days, 15:56:11.408681
train_1407	Finished	100%	0.58203125	2 days, 10:03:41.040259
train_1408	In Progress...	53.5%	0.6162109375	2 days, 15:56:11.408763
train_1409	Finished	100%	0.5517578125	2 days, 10:23:56.221825
train_1410	In Progress...	53.7%	0.4892578125	2 days, 15:56:11.408858
train_1411	Finished	100%	0.5029296875	2 days, 13:56:37.216733
train_1412	In Progress...	5.0%	0.3740234375	2 days, 15:56:11.408937
train_1413	In Progress...	9.9%	0.29296875	2 days, 15:56:11.409007
train_1414	In Progress...	4.7%	0.83984375	2 days, 15:56:11.409076
train_1415	In Progress...	3.2%	0.779296875	2 days, 15:56:11.409145

- Gave ray sessions to Vanessa on executing multiple experiments on DGX1.
- She is able to run the 10x experiments for case 2 on DGX1.

- We identified the issue with the bus error.
- Earlier, we assume that the issue is related to a possible shared memory limitation.
- This was because the 'free -mh' shows 1 TB of RAM on the nodes.
- In reality, SLURM limits the RAM, and this is the cause for Bus error.
- Simple solution: Ask for more RAM with srun/sbatch.

# Satori

---

- We verify that experiments are able to run on Satori nodes. Though pytorch version 1.3.
- On x86 machines, we use pytorch version 1.7
- Also, we should verify if the results are as expected to be completely sure for using Satori.



# Gitlab and Information Guide

---

- We create a final branch 'fixstars-dev' with all the scripts created.
- We also create an information guide to walkthrough the scripts and a quick parallelization process.